

COMPUTER ARCHITECTURE

Chapter 4: Microarchitecture



Introduction

- CPU performance factors
 - Instruction count
 - Determined by ISA and compiler
 - CPI and Cycle time
 - Determined by CPU hardware
- We will examine two MIPS implementations
 - A simplified version: $CPI = 1$
 - A more realistic pipelined version: $CPI \approx 1$
- Simple subset, shows most aspects
 - Memory reference: `lw`, `sw`
 - Arithmetic/logical: `add`, `sub`, `and`, `or`, `slt`
 - Control transfer: `beq`, `j`

Instruction execution

1. PC → **instruction memory (cache)**, fetch instruction
2. Register numbers → **registers file**, read registers
3. Depending on instruction class
 - 3.1. Use ALU to calculate
 - Arithmetic result → **done**
 - Memory address for load/store → **3.2**
 - Branch target address → **3.3**
 - 3.2. Access data memory for load/store
 - 3.3. PC ← target address or PC + 4


Execution stages

1. **Instruction fetch (IF)**: PC $\rightarrow$ instruction address
2. **Instruction decode (ID)**: register operands $\rightarrow$ register file
3. **Execute (EXE)**:
 - Load/store: compute a memory address
 - Arithmetic/logical: compute an arithmetic/logical result
4. **Memory access (MEM)**:
 - Load: read data memory
 - Store: write data memory
5. **Write back (WB)**:
 - Store a result of register file



Datapath - controller

- **Datapath**: contains information that is operated on by a functional unit
 - Instruction memory: contain instructions (**IF**)
 - Registers file: 32 32-bit registers (**ID** & **WB**)
 - ALU: calculate arithmetic and logical operations (**EXE**)
 - Data memory: contain data (**MEM**)
- **Control signals**: used for multiplexer selection or for directing the operation of a functional unit
 - Control unit
 - Multiplexer

Datapath overview




Multiplexer - MUX



$$S = 0 \rightarrow C = A$$


$$S = 1 \rightarrow C = B$$

Controller overview




Building datapath

- Hardware components:




Instruction fetch

- **Main operations:** fetching the next instruction from Instruction memory
 - **PC** → instruction address
 - **Instruction memory** → instruction (32 bits)
 - **PC** ← **PC** + 4 (using the **Add** component)
- **Results:**
 - 32 bits machine instruction (**Instruction**)
 - Address of the next instruction in PC



Instruction decode

- **Main operations:** Extract machine instructions
 - R-format, branch, and store: $rs, rt \rightarrow$ **Registers** $\rightarrow$ values
 - I-format: $rs \rightarrow$ **Registers** $\rightarrow$ values & $immediate \rightarrow$ **sign-extend**
- **Results:**
 - Values (32-bit) for the next stage



Instruction execution


- **Main operation**

- Calculate the arithmetic operations/
address of memory/register comparison

- R-format and **bne&beq**: both
operands collected from **Registers**
- I-format (except **bne&beq**): one
operands collected from Registers and
one from **sign-extend**

- **Results:**

- Arithmetic result/memory address (**ALU-
result**)
- Comparison result (**zero-output**)




Memory access

- **Main operations**

- Load: get data from **Data memory**
- Store: write data from **Registers** to **Data memory**


- **Results:**

- Load: values of **Data memory** (**Read data**)
- Store: N/A




Write back

- **Main operations:**
 - Write values back to **Registers** (arithmetic/load)
- **Result:**
 - N/A



Full datapath



Example

- **Question**: Assume that the processor is executing
$$\text{add } \$s0, \$s1, \$s2$$
 - Identify values of functional units' inputs/outputs
- **Answer**:
 - The machine code is: 000000_10001_10010_10000_00000_100000₂
 - Instruction memory:
 - **Instruction address** = PC (word address where we store the above instruction)
 - **Instruction** = machine code
 - Registers:
 - **Read register 1** = 10001₂ => **Read data 1** = content (\$s1)
 - **Read register 2** = 10010₂ => **Read data 2** = content (\$s2)
 - **Write register** = 10000₂ & **Write data** = content (\$s1) + content (\$s2)
 - ...sign-extend: input = 10000_00000_100000

Building controller


- Extracting bits from 32-bit instructions
 - Read register 1, Read register 2, and Write registers
 - Sign-extend
 - **Control** block
- Building a **Control** block:
 - Handling multiplexers
 - Handling control signals of functional units
 - *RegWrite*
 - *MemRead*
 - *MemWrite*
 - ...

Extracting bits

- **Registers:**
 - Read register 1 $\Leftarrow$ rs (instruction[25:21])
 - Read register 2 $\Leftarrow$ rt (instruction[20:16])
 - Write register $\Leftarrow$ rt/rd $\rightarrow$ need a multiplexer
- Sign-extend $\Leftarrow$ address (instruction[15:0])

R-type	0	rs	rt	rd	shamt	funct
	31:26	25:21	20:16	15:11	10:6	5:0
Load/ Store	35 or 43	rs	rt	address		
	31:26	25:21	20:16	15:0		
Branch	4	rs	rt	address		
	31:26	25:21	20:16	15:0		

Datapath with bit-selection



Control block


- Main function: handling multiplexers and functional units' control signals
 - ALU: two levels of decoding
 - Level 1: 6 bit **opcode** → 2 bit **ALUop**
 - Level 2: 2 bit **ALUop** (+ 6 bit function field) → 4 bit **ALU operation**
 - Muxes and control signals of functional units (**except ALU**): 6 bit **opcode**

- Predefined **ALU operation** values

Operator	ALU Opeation
and	0000
or	0001
add	0010
sub	0110
slt	0111

- **ALUop**:
 - opcode → add operator → **ALUop** = 00
 - opcode → sub operator → **ALUop** = 01
 - opcode → unknown → **ALUop** = 10

Full microarchitecture




Exercise


- **Question:** assume that registers store values of two times their numbers, for instance \$s1 stores values of $17 \times 2 = 34$, please identify:
 1. Which functional units contribute to the processing of following instructions
 2. Values of inputs/outputs of functional units when processing following instructions
 3. Values of control signals when processing following instructions
 - add \$t0, \$t1, \$t2
 - addi \$s0, \$s1, 100
 - lw \$s0, 100(\$s2) *# memory word at address 136 stores values of 2021*
 - sw \$s0, 100(\$s2)
 - beq \$s1, \$s0, L1

Unconditional jump instructions

- Chapter 2: update **PC** with concatenation of
 - Top 4 bits of old **PC**
 - 26-bit jump address
 - 00
$$PC = \{PC[31:28], \text{address}, 00\}$$
- One more way to update **PC** => Need a multiplexer & an extra control signal decoded from opcode



Datapath with Jumps added




Performance issue

- **Simplified version** (CPI = 1): every instruction executed in **only one cycle**
 - **Longest delay** determines clock period
 - What is the longest instruction?
 - Instruction memory → register file → ALU → data memory → register file
- Not feasible to vary period for different instructions
- We will improve performance by pipelining

Pipelining analogy

- Pipeline laundry:
 - Overlapping execution
 - Improving performance (time for entire group)
- Four loads
 - Speed-up = 2.3X
 - Not impressive
- Non-stop (#loads $\rightarrow \infty$)
 - Speed-up?
 - Number of stages



MIPS pipeline



- **Five stages**, one step per stage
 1. Instruction fetch (**IF**): PC → instruction address
 2. Instruction decode (**ID**): register operands → register file
 3. Execute (**EXE**):
 - Load/store: compute a memory address
 - Arithmetic/logical: compute an arithmetic/logical result
 4. Memory access (**MEM**):
 - Load: read data memory
 - Store: write data memory
 5. Write back (**WB**):
 - Store a result of register file

Pipeline performance

- Assume time for stages is
 - 100ps for register read or write (**ID** & **WB**)
 - 200ps for other stages (**IF**, **EXE**, & **MEM**)
- Compare **pipelined datapath** with **single-cycle datapath**

Instruction	Instr fetch	Register read	ALU op	Memory access	Register write	Total time
lw	200ps	100 ps	200ps	200ps	100 ps	800ps
sw	200ps	100 ps	200ps	200ps		700ps
R-format	200ps	100 ps	200ps		100 ps	600ps
beq	200ps	100 ps	200ps			500ps

Pipeline performance




Pipeline speed-up


$$\text{speed-up} = \frac{\text{time of single-cycle}}{\text{time of pipelined}} = \frac{\text{time bw instruction}_{\text{single-cycle}}}{\text{time bw instruction}_{\text{pipelined}}}$$

- If all stages are **balanced**:
 - speed_up = number of pipe stages
- If **not balanced**, speedup is less
- Source of speedup
 - **Throughput** increased
 - **Latency** (time for each instruction) does not decrease
 - Sometimes **increased**

Pipeline datapath




Instructions execution



Multi-Cycle Pipeline Diagram

How can we keep data when stages are not balanced?


Wholesale market example



- Trucks move forward when completed
 - **Accident** at the Apple store
 - **Barriers** can help: open every hour (**cycle**)

Pipeline registers

Barriers in wholesale markets $\Leftrightarrow$ Registers in digital circuits



Example


- **Question:** Given the following MIPS sequence:

1. lw \$s0, 20(\$s1)
2. sub \$t2, \$s2, \$s3
3. add \$t3, \$s3, \$s4
4. lw \$t4, 24(\$s1)
5. add \$t5, \$s5, \$s6


Assume that the sequence is executed by a 5-stage pipelined MIPS processor

- a) Draw a multi-cycle pipeline diagram for the sequence
- b) Analyze the 5th cycle with the datapath diagram in the previous slide

Multi-cycle pipeline diagram




Single-cycle pipeline diagram




Anything wrong?


Corrected pipeline datapath




Example of lw - IF




Example of lw - ID



Example of lw - EXE




Example of lw - MEM





Control signals

- Control signals travel across pipeline registers



Exercise


- **Question**: Given the following MIPS sequence:

```
lw    $s0, 20($s1)
sub    $t2, $s2, $s3
add    $t3, $s3, $s4
lw    $t4, 24($s1)
add    $t5, $s5, $s6
```

Assume that the sequence is executed by a 5-stage pipelined MIPS processor

- a) Identify values of control signals in cycle 5 **at the functional units**
- b) Identify values of control signals in cycle 5 **at the Control block**

Full pipeline micro-architecture



Hazards

- Situations that **prevent** starting the next instruction in the next cycle
- Structure hazard
 - A required **resource** is busy
- Data hazard
 - Need to wait for previous instruction to complete its **data** read/write
- Control hazard
 - Deciding on **control action** depends on previous instruction

Structure hazards


- **Conflict** for use of a resource
 - e.g.,: the laundry process, B forgot to bring clothes from the washing machine to dryer
- Should be eliminated entirely
 - **Stall** the pipe for that cycle
 - May repeat many times
- MIPS processors already **solved** all structure hazards
 - Separated instructions and data memory (caches)
 - Read and write registers use **different ports**

Data hazards

- An instruction **depends on completion of data** access by a previous instruction


```
sub $t0, $t2, $t3  
add $t3, $t0, $t1
```

- No any issues with the **simplified version**
- Problem** with pipeline



Control hazards

- **Branch** determines flow of control
 - Fetching next instruction depends on **branch outcome**
- Pipeline **updates PC** in the MEM stage
 - Still working on ID stage of branch



Data hazards solutions

1. Code rescheduling

- Done by **compiler** (software level)
- *Sometimes cannot find solutions*

2. Delay or stalls insertion


- Done by **hardware**
- **Always** can find solutions
- *Increase execution time & need extra hardware resources*

3. Forwarding

- Done by **hardware**
- *Cannot find solutions in a special case*
- Requires *extra hardware resources*

Code re-scheduling

- When **haven't** data hazards happened?
 - Read** after or in the same cycle with **Write** (**RAW**)



- Find** and **swap** **data-independent** instructions

Example


- **Question:** given the following sequence of MIPS instructions

1: lw \$t1, 0(\$t0)		1: lw \$t1, 0(\$t0)
2: lw \$t2, 4(\$t0)		2: lw \$t2, 4(\$t0)
3: add \$t3, \$t1, \$t2		5: lw \$t4, 8(\$t0)
4: sw \$t5, 12(\$t0)		7: addi \$t6, \$t0, 4
5: lw \$t4, 8(\$t0)		3: add \$t3, \$t1, \$t2
6: add \$t5, \$t1, \$t4		4: sw \$t5, 12(\$t0)
7: addi \$t6, \$t0, 4		6: add \$t5, \$t1, \$t4

- Identify data hazards and solve by code re-scheduling
- **Answer:**
 - The third and the sixth instructions (add \$t3, \$t1, \$t2 and add \$t5, \$t1, \$t4) have data hazards
 - Move two data-independent instructions 5 and 7 to right after the second instruction

Stalls insertion

- When **haven't** data hazards happened?
 - Read** after or in the same cycle with **Write** (**RAW**)



- Delay** instructions that use data:
 - ID** stage of the instruction **using** data $\equiv$ **WB** stage of the instruction **producing** data

Example

- **Questions:** given the following MIPS sequence of instructions

1: lw \$t1, 0(\$t0)


2: lw \$t2, 4(\$t0)

3: add \$t3, \$t1, \$t2

4: sw \$t3, 12(\$t0)


- Identify data hazards and solve them by the stalls insertion method; how many cycles needed for the sequence?
- **Answer:**
 - Data hazards (2) - (3) & (3) - (4)
 - Insert two stalls for the 3rd instruction & two stalls for the 4th instruction
 - 12 cycles needed

Example (cont.)



Forwarding

- When can an instruction use data at **soonest**?
 - **Right after** data is **produced** (**EXE** stage & **MEM** stage)
- When is data **processed**?
 - **Apply** operators: **EXE** state



- **Use** data right after created

Example


- **Question**: given the following MIPS sequence

```
1: sub $s2, $s1, $s3
2: and $s7, $s2, $s5
3: or $s8, $s6, $s2
4: add $s0, $s2, $s2
5: sw $s5, 100($s2)
```

- Analyze data dependencies & identify data hazards
- **Answer**:
 - \$s2 **produced** by the 1st instruction **used** by all other instructions
 - **Data hazards**: 2nd & 3rd instructions


Example (cont.)

- Answer (cont.):




Forwarding

- Consider the previous sequence of MIPS instructions
 - Data can be **forwarded**



Load used data hazards

- When instruction producing data is a **load**, can data be forwarded?
 - Cannot forward** to the next instruction since data is produced **later**




- Delay:**
 - 1 stall** with forwarding or **2 stalls** without forwarding

Detecting data hazards

- **EXE** hazard **versus** **MEM** hazard
 - **EXE** hazard: forward from the **EX/MEM** register
 - **Destination register** in the MEM stage $\equiv$ one of the source registers of the instruction in EXE
 - **MEM** hazard: forward from the **MEM/WB** register
 - **Destination register** in the WB stage $\equiv$ one of the source registers of the instruction in EXE
- **Passing** register numbers along the pipe
 - ID/EX.rs & ID/EX.rt: first & second sources
 - EX/MEM.rd; MEM/WB.rd: destinations

EX hazard conditions



1a. $ID/EX.rs = EX/MEM.rd$

1b. $ID/EX.rt = EX/MEM.rd$ (not happened when an **I-format** instruction is in the **EXE** stage)

MEM hazard conditions

- Consider the following MIPS sequence of instructions

```
1: add $s0, $s0, $s1
2: add $s0, $s0, $s2
3: add $s0, $s0, $s3
```


- Both **EX** and **MEM hazards** seem occur
 - EX hazard is **correct**
- MEM hazard conditions

2a. $(ID/EX.rs = MEM/WB.rd) \ \& \ (ID/EX.rs \neq EX/MEM.rd)$

2b. $(ID/EX.rt = MEM/WB.rd) \ \& \ (ID/EX.rt \neq EX/MEM.rd)$

(not happened when an **I-format** instruction is in the **EXE** stage)

Datapath with forwarding



Forwarding control values

Mux values (binary)	Source	Explanation
F1 = 00	ID/EX	The first ALU operand comes from the registers file
F1 = 10	EX/MEM	The first ALU operand is forwarded from the prior ALU result
F1 = 01	MEM/WB	The first ALU operand is forwarded from data memory of an earlier ALU result
F2 = 00	ID/EX	The second ALU operand comes from the registers file
F2 = 10	EX/MEM	The second ALU operand is forwarded from the prior ALU result
F2 = 01	MEM/WB	The second ALU operand is forwarded from data memory of an earlier ALU result

Forwarding conditions

- EX hazard


- 1a: if (EX/MEM.RegWrite and (EX/MEM.rd $\neq$ 0) and (EX/MEM.rd = ID/EX.rs)) **F1 = 10**
- 1b: if (EX/MEM.RegWrite and (EX/MEM.rd $\neq$ 0) and (EX/MEM.rd = ID/EX.rt)) **F2 = 10**

- MEM hazard

- 2a: if (MEM/WB.RegWrite and (MEM/WB.rd $\neq$ 0) and not (EX/MEM.RegWrite and (EX/MEM.rd $\neq$ 0) and (EX/MEM.rd = ID/EX.rs)) and (MEM/WB.rd = ID/EX.rs)) **F1 = 01**
- 2b: if (MEM/WB.RegWrite and (MEM/WB.rd $\neq$ 0) and not (EX/MEM.RegWrite and (EX/MEM.rd $\neq$ 0) and (EX/MEM.rd = ID/EX.rt)) and (MEM/WB.rd = ID/EX.rt)) **F2 = 01**


Detecting Load-use data hazards

- Check when using data instruction is decoded in **ID** stage
 - Producing data instruction (**load**) is in the **EXE** stage
 - The **sooner** the **better** due to a stall inserted
- Load-use hazard when
 - **ID/EX.MemRead** and $((\text{ID/EX.r}_t = \text{IF/ID.rs}) \text{ or } (\text{ID/EX.r}_t = \text{IF/ID.r}_t))$
- **How to stall the pipeline?**
 - **Force** control signals in **ID/EX** register to 0 $\Rightarrow$ EXE, MEM, and WB do nothing
 - **Prevent** updating PC and **ID/EX** register



Branch hazards solutions

- **Predict** outcome of branch
 - Only **stall** if prediction is wrong



Branch prediction

- **Static** branch prediction
 - Based on **typical branch behavior**
 - Example: loop and if-statement branches
 - Predict backward branches taken
 - Predict forward branches not taken
- **Dynamic** branch prediction
 - Hardware measures **actual branch behavior**
 - e.g., record recent history of each branch
 - Assume future behavior will continue the trend
 - When wrong, stall while re-fetching, and **update history**

Concluding remarks

- ISA influences design of **datapath** and **control**
- Datapath and control influence design of ISA
- **Pipelining** improves instruction throughput using parallelism
 - More instructions completed per second
 - Latency for each instruction not reduced
- **Hazards**: structural, data, control

The end

